

Sri Krishna Loganathan

Dallas, TX (open to relocate) | srikrishnaloganathan@gmail.com | +1 (945)-946-1931 | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

A Data engineer with 3+ years of industry experience in applying Data Modeling techniques to build efficient Data Architecture and ETL designs, leveraging Data Warehouse methodologies to deliver scalable, high-quality analytical systems. Ability to turn complex data into clear & trusted insights empowering leadership to make confident decisions and enable revenue growth.

SKILLS

- Languages & Libraries: Python (Pandas, NumPy, PySpark, Spark SQL), SQL, PL/SQL
 - Data Engineering & Cloud: GCP, Dataflow, Pub/Sub, Cloud Run, Schedulers (Apache Airflow, Composer, Cron), Databricks, GCS, Azure DevOps, Docker, Hadoop, RESTful APIs
 - Databases & Warehouses: BigQuery, Snowflake, MSOL, SSMS, PostgreSQL, Hive
 - Visualization & BI: Tableau, Power BI, Microsoft Excel, Looker
 - Domain knowledge: OOP's, Agile SDLC, JIRA, RAG, Agentic AI, GenAI
-

PROFESSIONAL EXPERIENCE

Conifer Health Solutions | Data Engineer

January 2026 – May 2026

- Built a Change Data Capture (CDC)-driven ETL pipeline in BigQuery tracking 50K+ monthly claim transitions, enabling automated reconciliation and production deployments via Azure DevOps CI/CD workflows and Git-based release pipelines
- Migrated 50+ Alteryx workflows to BigQuery by integrating multi-source data (IBM DB2, SQL Server, GCS) and reconstructing complex workflow logic into optimized cloud SQL, improving data processing scalability
- Reverse-engineered 20+ client reporting data marts to establish column-level data lineage and redesigned the data model into a Star Schema with a conformed dimension supporting 5 dashboards, reducing dependencies and improving report reliability

Tenet Healthcare | Data Engineer

June 2025 – December 2025

- Enhanced orchestration framework by introducing an execution order parameter in bundle-driven Airflow DAGs, enabling 30+ stored procedures to run in parallel and reducing end-to-end workflow runtime by 60%
- Redesigned dimensional model by implementing Type 1 SCD merge logic and decoupling fact and dimension tables using hash keys, enabling parallelized batch loads in BigQuery and reducing processing time by 70%
- Re-engineered 150+ legacy stored procedures into BigQuery optimized SQL leveraging partitioned/clustered tables and ephemeral staging constructs, while automating reconciliation and integration test suites through Airflow DAGs in Cloud Composer, achieving 100% validation accuracy and reducing manual QA effort by 40%
- Developed a metrics-based Data Validation Testing (DVT) framework using Python and SQL to automate post-migration validation between legacy systems and BigQuery, ensuring data accuracy and reliability for BI and analytics teams

Alight Solutions - contract via Mphasis Limited | Technical Systems Analyst

August 2022 – May 2024

- Identified 5+ critical production outage in enterprise KPI reporting by analyzing support tickets and AWS infrastructure, performing root cause analysis that uncovered S3 policy misconfigurations blocking data ingestion pipelines
 - Identified S3 bucket access issues and RDS schema mismatches affecting data pipelines by analyzing AWS CloudWatch logs; led incident resolution and SLA enforcement with infrastructure teams, achieving 40% MTTR reduction, elimination of ~150 monthly KPI reconciliation errors, and prevention of 100+ production failures
 - Proactively led cross-functional escalation initiatives alongside senior management and technical teams to identify and mitigate 100+ potential outages, authoring 10+ knowledge bases and driving a 40% NPS improvement through structured communication and strengthened ITSM practices using Service Now
-

TECHNICAL PROJECTS

Formula 1 Data Lakehouse for ML Pipelines | Databricks, PySpark & Tableau

February 2026 – May 2026

- Engineered a metadata-driven bronze ingestion framework on Databricks using PySpark to ingest 11 raw F1 datasets from Google Drive into Unity Catalog Volumes as Delta Lake tables, automating schema validation and enabling RAG pipelines to retrieve contextual race data with 95%+ accuracy
- Designed a Medallion architecture with star schema on Databricks, consolidating 11 Bronze tables into a curated Silver layer (5 dimensions, 1 fact table, 35+ analytical measures), enabling real-time performance dashboards and NL2SQL AI pipelines for predictive race analytics

Agentic Insurance Claims System | LangGraph, LLM & Streamlit

September 2025 – December 2025

- Architected multi-agent LLM system (LangGraph) with specialized roles (policy validator, fraud detector, settlement scorer) using PostgreSQL for state persistence and Gemini Vision API for multimodal analysis. Implemented dynamic agent routing with RAG context injection (3-claim timeline), achieving 94% approval accuracy and 89% fraud detection precision
 - Built Streamlit platform orchestrating sequential agentic workflows combining coverage verification, image forensics, and claims scoring. Automated 85% of manual review workload while processing 100+ monthly claims with explainable decision traces
-

EDUCATION

The University of Texas at Dallas | Master's, Information Technology (GPA: 3.6/4.0)

August 2024 – May 2026

B.M.S College of Engineering | Bachelor's, Telecommunication Engineering (GPA: 3.5/4.0)

August 2018 – May 2022